

ISMT S-136 Time Series Analysis with Python

Harvard Summer School

Dmitry Kurochkin

Summer 2021

Lecture 11

Contents

1 Recurrent Neural Networks (RNN)

- Recurrent Neuron and Layer of Recurrent Neurons
- Memory Cells
- Long Short-Term Memory (LSTM) Cell
- Gated Recurrent Unit (GRU) Cell

2 Unstable Gradients

- Vanishing/Exploding Gradients Problems
- Techniques to Alleviate the Unstable Gradient Problems

3 Neural Network Optimization Algorithms

- Objective (Cost) Function
- SGD, mini-batch GD, and GD Optimization
- Forward Propagation and Backpropagation
- Momentum Optimization, NAG, AdaGrad, RMSProp, and Adam

Contents

1 Recurrent Neural Networks (RNN)

- Recurrent Neuron and Layer of Recurrent Neurons
- Memory Cells
- Long Short-Term Memory (LSTM) Cell
- Gated Recurrent Unit (GRU) Cell

2 Unstable Gradients


- Vanishing/Exploding Gradients Problems
- Techniques to Alleviate the Unstable Gradient Problems

3 Neural Network Optimization Algorithms

- Objective (Cost) Function
- SGD, mini-batch GD, and GD Optimization
- Forward Propagation and Backpropagation
- Momentum Optimization, NAG, AdaGrad, RMSProp, and Adam


RNN: Recurrent Neuron

Recurrent Neuron:



RNN: Layer of Recurrent Neurons

Layer of Recurrent Neurons:



Layer of Recurrent Neurons: Keras

Simple RNN in Keras:

```
n_features = 2
n_timesteps = 200

model = models.Sequential()
model.add(layers.SimpleRNN(3, activation='relu', input_shape=(n_timesteps,n_features)))
model.add(layers.Dense(1, activation='linear'))

model.summary()
```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
simple_rnn_6 (SimpleRNN)	(None, 3)	18
dense_6 (Dense)	(None, 1)	4

Total params: 22

Trainable params: 22

Non-trainable params: 0

Contents

1 Recurrent Neural Networks (RNN)

- Recurrent Neuron and Layer of Recurrent Neurons
- **Memory Cells**
- Long Short-Term Memory (LSTM) Cell
- Gated Recurrent Unit (GRU) Cell

2 Unstable Gradients


- Vanishing/Exploding Gradients Problems
- Techniques to Alleviate the Unstable Gradient Problems

3 Neural Network Optimization Algorithms

- Objective (Cost) Function
- SGD, mini-batch GD, and GD Optimization
- Forward Propagation and Backpropagation
- Momentum Optimization, NAG, AdaGrad, RMSProp, and Adam

RNN: Memory Cells

Layer of Recurrent Neurons:



Contents

1 Recurrent Neural Networks (RNN)

- Recurrent Neuron and Layer of Recurrent Neurons
- Memory Cells
- Long Short-Term Memory (LSTM) Cell
- Gated Recurrent Unit (GRU) Cell

2 Unstable Gradients


- Vanishing/Exploding Gradients Problems
- Techniques to Alleviate the Unstable Gradient Problems

3 Neural Network Optimization Algorithms

- Objective (Cost) Function
- SGD, mini-batch GD, and GD Optimization
- Forward Propagation and Backpropagation
- Momentum Optimization, NAG, AdaGrad, RMSProp, and Adam

Long Short-Term Memory (LSTM) Cell

LSTM Cell:



Long Short-Term Memory (LSTM) Cell

LSTM Cell Model:

$$\begin{aligned}\mathbf{i}_{(t)} &= \sigma(\mathbf{W}_{xi}^T \cdot \mathbf{x}_{(t)} + \mathbf{W}_{hi}^T \cdot \mathbf{h}_{(t-1)} + \mathbf{b}_i) \\ \mathbf{f}_{(t)} &= \sigma(\mathbf{W}_{xf}^T \cdot \mathbf{x}_{(t)} + \mathbf{W}_{hf}^T \cdot \mathbf{h}_{(t-1)} + \mathbf{b}_f) \\ \mathbf{o}_{(t)} &= \sigma(\mathbf{W}_{xo}^T \cdot \mathbf{x}_{(t)} + \mathbf{W}_{ho}^T \cdot \mathbf{h}_{(t-1)} + \mathbf{b}_o) \\ \mathbf{g}_{(t)} &= \tanh(\mathbf{W}_{xg}^T \cdot \mathbf{x}_{(t)} + \mathbf{W}_{hg}^T \cdot \mathbf{h}_{(t-1)} + \mathbf{b}_g) \\ \mathbf{c}_{(t)} &= \mathbf{f}_{(t)} \otimes \mathbf{c}_{(t-1)} + \mathbf{i}_{(t)} \otimes \mathbf{g}_{(t)} \\ \mathbf{y}_{(t)} &= \mathbf{h}_{(t)} = \mathbf{o}_{(t)} \otimes \tanh(\mathbf{c}_{(t)})\end{aligned}$$

- \mathbf{W}_{xi} , \mathbf{W}_{xf} , \mathbf{W}_{xo} , \mathbf{W}_{xg} are the weight matrices of each of the four layers for their connection to the input vector $\mathbf{x}_{(t)}$.
- \mathbf{W}_{hi} , \mathbf{W}_{hf} , \mathbf{W}_{ho} , and \mathbf{W}_{hg} are the weight matrices of each of the four layers for their connection to the previous short-term state $\mathbf{h}_{(t-1)}$.
- \mathbf{b}_i , \mathbf{b}_f , \mathbf{b}_o , and \mathbf{b}_g are the bias terms for each of the four layers. Note that TensorFlow initializes \mathbf{b}_f to a vector full of 1s instead of 0s. This prevents forgetting everything at the beginning of training.

Applications of LSTM Cells

- greatly improved speech recognition on over 4 billion Android phones (since mid 2015)
- greatly improved machine translation through Google Translate (since Nov 2016)
- greatly improved machine translation through Facebook (over 4 billion LSTMbased translations per day as of 2017)
- Siri and Quicktype on almost 2 billion iPhones (since 2016)
- generating answers by Amazon's Alexa and numerous other similar applications.

Long Short-Term Memory (LSTM) Cell: Keras

LSTM in Keras:

```
n_features = 2
n_timesteps = 200

model = models.Sequential()
model.add(LSTM(16, activation='relu', input_shape=(n_timesteps,n_features)))
model.add(layers.Dense(1, activation='linear'))

model.summary()
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
=====	=====	=====
lstm_1 (LSTM)	(None, 16)	1216
dense_7 (Dense)	(None, 1)	17
=====	=====	=====

Total params: 1,233

Trainable params: 1,233

Non-trainable params: 0

Contents

1 Recurrent Neural Networks (RNN)

- Recurrent Neuron and Layer of Recurrent Neurons
- Memory Cells
- Long Short-Term Memory (LSTM) Cell
- Gated Recurrent Unit (GRU) Cell

2 Unstable Gradients

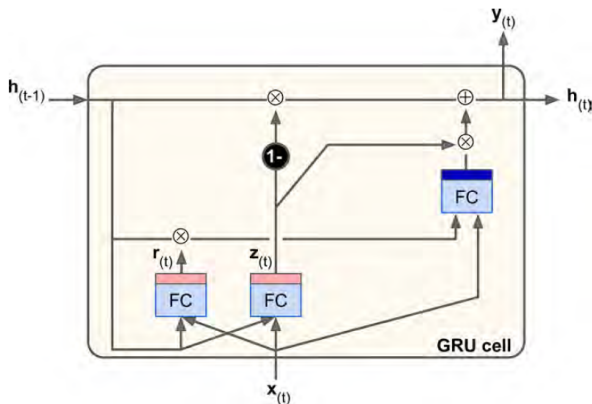
- Vanishing/Exploding Gradients Problems
- Techniques to Alleviate the Unstable Gradient Problems

3 Neural Network Optimization Algorithms

- Objective (Cost) Function
- SGD, mini-batch GD, and GD Optimization
- Forward Propagation and Backpropagation
- Momentum Optimization, NAG, AdaGrad, RMSProp, and Adam

Gated Recurrent Unit (GRU) Cell

GRU Cell:



Gated Recurrent Unit (GRU) Cell

GRU Cell Model:

$$\begin{aligned} \mathbf{z}_{(t)} &= \sigma(\mathbf{W}_{xz}^T \cdot \mathbf{x}_{(t)} + \mathbf{W}_{hz}^T \cdot \mathbf{h}_{(t-1)}) \\ \mathbf{r}_{(t)} &= \sigma(\mathbf{W}_{xr}^T \cdot \mathbf{x}_{(t)} + \mathbf{W}_{hr}^T \cdot \mathbf{h}_{(t-1)}) \\ \mathbf{g}_{(t)} &= \tanh(\mathbf{W}_{xg}^T \cdot \mathbf{x}_{(t)} + \mathbf{W}_{hg}^T \cdot (\mathbf{r}_{(t)} \otimes \mathbf{h}_{(t-1)})) \\ \mathbf{h}_{(t)} &= (1 - \mathbf{z}_{(t)}) \otimes \tanh(\mathbf{W}_{xg}^T \cdot \mathbf{h}_{(t-1)} + \mathbf{z}_{(t)} \otimes \mathbf{g}_{(t)}) \end{aligned}$$

Contents

1 Recurrent Neural Networks (RNN)

- Recurrent Neuron and Layer of Recurrent Neurons
- Memory Cells
- Long Short-Term Memory (LSTM) Cell
- Gated Recurrent Unit (GRU) Cell

2 Unstable Gradients

- Vanishing/Exploding Gradients Problems
- Techniques to Alleviate the Unstable Gradient Problems

3 Neural Network Optimization Algorithms


- Objective (Cost) Function
- SGD, mini-batch GD, and GD Optimization
- Forward Propagation and Backpropagation
- Momentum Optimization, NAG, AdaGrad, RMSProp, and Adam

Vanishing/Exploding Gradients Problems

Unstable gradients:

- 1 *Vanishing gradients* problem: Given current weights w of the NN and inputs (data), the gradient of the activation function may be very small resulting in the corresponding weights virtually unchanged during the iterations / updates.
- 2 *Exploding gradients* problem: The weights may on the contrary blow up - this problem is mostly encountered in recurrent neural networks.

Sigmoid function:



Contents

1 Recurrent Neural Networks (RNN)

- Recurrent Neuron and Layer of Recurrent Neurons
- Memory Cells
- Long Short-Term Memory (LSTM) Cell
- Gated Recurrent Unit (GRU) Cell

2 Unstable Gradients

- Vanishing/Exploding Gradients Problems
- Techniques to Alleviate the Unstable Gradient Problems

3 Neural Network Optimization Algorithms

- Objective (Cost) Function
- SGD, mini-batch GD, and GD Optimization
- Forward Propagation and Backpropagation
- Momentum Optimization, NAG, AdaGrad, RMSProp, and Adam

Techniques to Alleviate the Unstable Gradient Problems

Ways to resolve the problems:

- 1 “Proper” initialization of weights: special initial distribution, reusing pretrained layers, etc.
- 2 Nonsaturating activations functions: Leaky ReLU, exponential LU (ELU), etc.
- 3 Batch normalization (BN): scale inputs before each layer during training (two more parameters)
- 4 Gradient clipping: set a threshold for the gradient

Contents

1 Recurrent Neural Networks (RNN)

- Recurrent Neuron and Layer of Recurrent Neurons
- Memory Cells
- Long Short-Term Memory (LSTM) Cell
- Gated Recurrent Unit (GRU) Cell

2 Unstable Gradients

- Vanishing/Exploding Gradients Problems
- Techniques to Alleviate the Unstable Gradient Problems

3 Neural Network Optimization Algorithms

- **Objective (Cost) Function**
- SGD, mini-batch GD, and GD Optimization
- Forward Propagation and Backpropagation
- Momentum Optimization, NAG, AdaGrad, RMSProp, and Adam

Objective (Cost) Function

Suppose we want to train a supervised model (e.g., Neural Network) using a set of observations:

$$(\mathbf{x}_1, \mathbf{y}_1), (\mathbf{x}_2, \mathbf{y}_2), (\mathbf{x}_3, \mathbf{y}_3), \dots, (\mathbf{x}_m, \mathbf{y}_m)$$

then we define the objective (or cost) function as mean loss:

$$J(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m L^{(i)}(\mathbf{w}),$$

where

$$L^{(i)}(\mathbf{w}) = L(\underbrace{\hat{\mathbf{y}}^{(i)}(\mathbf{w})}_{\text{prediction}}, \underbrace{\mathbf{y}^{(i)}}_{\text{observed}})$$

is the loss associated with a single observation i .

Objective (Cost) Function

The list of the most common cost functions:

- Mean Squared Error:

$$J(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m \sum_{j=1}^M (\hat{y}_j^{(i)} - y_j^{(i)})^2$$

- Mean Absolute Error:

$$J(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m \left(\sum_{j=1}^M (\hat{y}_j^{(i)} - y_j^{(i)})^2 \right)^{\frac{1}{2}}$$

- Cross-Entropy:

$$J(\mathbf{w}) = -\frac{1}{m} \sum_{i=1}^m \sum_{j=1}^M y_j^{(i)} \ln \hat{y}_j^{(i)}$$

Contents

1 Recurrent Neural Networks (RNN)

- Recurrent Neuron and Layer of Recurrent Neurons
- Memory Cells
- Long Short-Term Memory (LSTM) Cell
- Gated Recurrent Unit (GRU) Cell

2 Unstable Gradients

- Vanishing/Exploding Gradients Problems
- Techniques to Alleviate the Unstable Gradient Problems

3 Neural Network Optimization Algorithms

- Objective (Cost) Function
- **SGD, mini-batch GD, and GD Optimization**
- Forward Propagation and Backpropagation
- Momentum Optimization, NAG, AdaGrad, RMSProp, and Adam

SGD, mini-batch GD, and GD Optimization: 'sgd'

The SGD, mini-batch GD, and GD Optimization (with learning rate α) are all defined as follows:

$$\mathbf{w} := \mathbf{w} - \alpha \underbrace{\frac{1}{s} \sum_{i=1}^s \nabla L^{(i)}(\mathbf{w})}_{\approx \nabla J(\mathbf{w})},$$

where $L^{(i)}(\mathbf{w})$ is based on one observation i and

- $s = 1$ in case of Stochastic Gradient Descent (SGD)
- $1 < s < m$ in case of mini-batch Gradient Descent (mini-batch GD)
- $s = m$ in case of Gradient Descent (GD)

Here, m denotes the total number of observations in the data set.

SGD, mini-batch GD, and GD Optimization

Example: Mini-batch GD with $s = 128$ and $\alpha = 0.01$.

```
model = models.Sequential()
model.add(layers.Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(layers.Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(layers.Dense(10, activation='softmax'))

model.summary()
```

Model: "sequential_14"

Layer (type)	Output Shape	Param #
dense_35 (Dense)	(None, 512)	401920
dropout_1 (Dropout)	(None, 512)	0
dense_36 (Dense)	(None, 512)	262656
dropout_2 (Dropout)	(None, 512)	0
dense_37 (Dense)	(None, 10)	5130
Total params: 669,706		
Trainable params: 669,706		
Non-trainable params: 0		

```
nepochs = 35
model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='sgd')

history = model.fit(X_train, y_train,
                    batch_size=128, epochs=nepochs,
                    verbose=1,
                    validation_data=(X_test, y_test))
```

SGD, mini-batch GD, and GD Optimization

Example: Mini-batch GD with $s = 128$ and $\alpha = 0.05$.

```
model = models.Sequential()
model.add(layers.Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(layers.Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(layers.Dense(10, activation='softmax'))

model.summary()
```

Model: "sequential_14"

Layer (type)	Output Shape	Param #
dense_35 (Dense)	(None, 512)	401920
dropout_1 (Dropout)	(None, 512)	0
dense_36 (Dense)	(None, 512)	262656
dropout_2 (Dropout)	(None, 512)	0
dense_37 (Dense)	(None, 10)	5130
Total params: 669,706		
Trainable params: 669,706		
Non-trainable params: 0		

```
nepochs = 35
model.compile(loss='categorical_crossentropy', metrics=['accuracy'],
              optimizer=keras.optimizers.SGD(lr=0.05))

history = model.fit(X_train, y_train,
                    batch_size=128, epochs=nepochs,
                    verbose=1,
                    validation_data=(X_test, y_test))
```

Contents

1 Recurrent Neural Networks (RNN)

- Recurrent Neuron and Layer of Recurrent Neurons
- Memory Cells
- Long Short-Term Memory (LSTM) Cell
- Gated Recurrent Unit (GRU) Cell

2 Unstable Gradients

- Vanishing/Exploding Gradients Problems
- Techniques to Alleviate the Unstable Gradient Problems

3 Neural Network Optimization Algorithms

- Objective (Cost) Function
- SGD, mini-batch GD, and GD Optimization
- **Forward Propagation and Backpropagation**
- Momentum Optimization, NAG, AdaGrad, RMSProp, and Adam

Forward Propagation

Let's again consider a Neural Network (NN) with

n inputs,

M outputs and

1 hidden layer with H neurons.

The Forward Propagation is then

$$\hat{y}_m = \sigma_m^{(2)} \left(\sum_{h=0}^H w_{hm}^{(2)} \underbrace{\sigma_h^{(1)} \left(\sum_{j=0}^n w_{jh}^{(1)} x_j \right)}_{\doteq u_h} \right).$$

Backpropagation

Given an observation (\mathbf{x}, \mathbf{y}) , assume we want to minimize the Mean Squared Error Loss

$$L(\mathbf{w}) = \sum_{m=1}^M (\hat{y}_m - y_m)^2,$$

where

$$\hat{y}_m = \sigma_m^{(2)} \left(\sum_{h=0}^H w_{hm}^{(2)} \underbrace{\sigma_h^{(1)} \left(\sum_{j=0}^n w_{jh}^{(1)} x_j \right)}_{\doteq \mathbf{u}_h} \right).$$

Then need to compute $\frac{L(\mathbf{w})}{\partial w_{hm}^{(2)}}$ and $\frac{L(\mathbf{w})}{\partial w_{jh}^{(1)}}$.

But we know derivatives of $\sigma_m^{(2)}$ and $\sigma_h^{(1)}$ exactly!

Also, we have $\sum_{j=0}^n w_{jh}^{(1)} x_j$, \mathbf{u}_h , $\sum_{h=0}^H w_{hm}^{(2)} \mathbf{u}_h$, and \hat{y}_m computed during forward propagation!

Example of Forward Propagation/Backpropagation

Let's consider the following Neural Network:

input layer	hidden layer	output layer
x_1	$u_1 = f(\underbrace{w_{01}^{(1)} + w_{11}^{(1)}x_1 + w_{21}^{(1)}x_2}_{z_1^{(1)}})$	$\hat{y} = f(\underbrace{w_0^{(2)} + w_1^{(2)}u_1 + w_2^{(2)}u_2}_{z^{(2)}})$
x_2	$u_2 = f(\underbrace{w_{02}^{(1)} + w_{12}^{(1)}x_1 + w_{22}^{(1)}x_2}_{z_2^{(1)}})$	

Here, $f(x)$ denotes the activation function, for example, ReLU.

Example of Forward Propagation/Backpropagation

Let's consider the following Neural Network:

input layer	hidden layer	output layer
x_1	$u_1 = f(\underbrace{w_{01}^{(1)} + w_{11}^{(1)}x_1 + w_{21}^{(1)}x_2}_{z_1^{(1)}})$	$\hat{y} = f(\underbrace{w_0^{(2)} + w_1^{(2)}u_1 + w_2^{(2)}u_2}_{z^{(2)}})$
x_2	$u_2 = f(\underbrace{w_{02}^{(1)} + w_{12}^{(1)}x_1 + w_{22}^{(1)}x_2}_{z_2^{(1)}})$	

Here, $f(x)$ denotes the activation function, for example, ReLU.

Forward Propagation: Given weights w and inputs x_1, x_2 , compute

- $z_1^{(1)}$ and $z_2^{(1)}$
- u_1 and u_2
- \hat{y}

Example of Forward Propagation/Backpropagation

Backpropagation:

Given weights w , inputs x_1, x_2 , and $z_1^{(1)}, z_2^{(1)}, u_1, u_2, \hat{y}$, compute

- Error associated with the output layer:

$$\varepsilon^{(2)} \doteq \frac{\partial L}{\partial \hat{y}} = \frac{\partial}{\partial \hat{y}} [(\hat{y} - y)^2] = 2(\hat{y} - y)$$

- Errors associated with the hidden layer:

$$\varepsilon_h^{(1)} \doteq \frac{\partial L}{\partial u_h} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u_h} = \varepsilon^{(2)} f'(z^{(2)}) w_h^{(2)}, \quad h = 1, 2.$$

Example of Forward Propagation/Backpropagation

Computation of $\nabla L(\mathbf{w})$:

- Partial derivatives of the loss function with respect to weights in the output layer:

$$\frac{\partial L}{\partial w_h^{(2)}} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial w_h^{(2)}} = \varepsilon^{(2)} \frac{\partial}{\partial w_h^{(2)}} \left[f(\underbrace{w_0^{(2)} + w_1^{(2)} u_1 + w_2^{(2)} u_2}_{z^{(2)}}) \right] = \varepsilon^{(2)} f'(z^{(2)}) u_h,$$

where $h = 0, 1, 2$.

- Partial derivatives of the loss function with respect to weights in the hidden layer:

$$\frac{\partial L}{\partial w_{jh}^{(1)}} = \frac{\partial L}{\partial u_h} \frac{\partial u_h}{\partial w_{jh}^{(1)}} = \varepsilon_h^{(1)} \frac{\partial}{\partial w_{jh}^{(1)}} \left[f(\underbrace{w_{0h}^{(1)} + w_{1h}^{(1)} x_1 + w_{2h}^{(1)} x_2}_{z_h^{(1)}}) \right] = \varepsilon_h^{(1)} f'(z_h^{(1)}) x_j,$$

for each $j = 0, 1, 2$ and $h = 1, 2$. Here, we define $x_0 \doteq 1$.

Example of Forward Propagation/Backpropagation

The Stochastic Gradient Descent (SGD) update of the weights using learning rate α :

$$\mathbf{w} := \mathbf{w} - \alpha \nabla L,$$

where $\nabla L \doteq \left(\underbrace{\frac{\partial L}{\partial w_{01}^{(1)}}, \frac{\partial L}{\partial w_{11}^{(1)}}, \frac{\partial L}{\partial w_{21}^{(1)}}, \frac{\partial L}{\partial w_{02}^{(1)}}, \frac{\partial L}{\partial w_{12}^{(1)}}, \frac{\partial L}{\partial w_{22}^{(1)}}}_{\text{hidden layer}}, \underbrace{\frac{\partial L}{\partial w_0^{(2)}}, \frac{\partial L}{\partial w_1^{(2)}}, \frac{\partial L}{\partial w_2^{(2)}}}_{\text{output layer}} \right)^T$.

Therefore,

$$\begin{aligned} \mathbf{w} &:= \mathbf{w} - \alpha \nabla L \\ &= \left(\underbrace{w_{01}^{(1)}, w_{11}^{(1)}, w_{21}^{(1)}, w_{02}^{(1)}, w_{12}^{(1)}, w_{22}^{(1)}}_{\text{hidden layer}}, \underbrace{w_0^{(2)}, w_1^{(2)}, w_2^{(2)}}_{\text{output layer}} \right)^T \\ &\quad - \alpha \left(\underbrace{\frac{\partial L}{\partial w_{01}^{(1)}}, \frac{\partial L}{\partial w_{11}^{(1)}}, \frac{\partial L}{\partial w_{21}^{(1)}}, \frac{\partial L}{\partial w_{02}^{(1)}}, \frac{\partial L}{\partial w_{12}^{(1)}}, \frac{\partial L}{\partial w_{22}^{(1)}}}_{\text{hidden layer}}, \underbrace{\frac{\partial L}{\partial w_0^{(2)}}, \frac{\partial L}{\partial w_1^{(2)}}, \frac{\partial L}{\partial w_2^{(2)}}}_{\text{output layer}} \right)^T \end{aligned}$$

Contents

1 Recurrent Neural Networks (RNN)

- Recurrent Neuron and Layer of Recurrent Neurons
- Memory Cells
- Long Short-Term Memory (LSTM) Cell
- Gated Recurrent Unit (GRU) Cell

2 Unstable Gradients

- Vanishing/Exploding Gradients Problems
- Techniques to Alleviate the Unstable Gradient Problems

3 Neural Network Optimization Algorithms

- Objective (Cost) Function
- SGD, mini-batch GD, and GD Optimization
- Forward Propagation and Backpropagation
- Momentum Optimization, NAG, AdaGrad, RMSProp, and Adam

Momentum Optimization

The Momentum Optimization algorithm is defined as follows:

First, *momentum vector* \mathbf{v} is initialized at $\mathbf{0}$ and then the updates are

$$\mathbf{v} := -\alpha \underbrace{\frac{1}{s} \sum_{i=1}^s \nabla L^{(i)}(\mathbf{w})}_{\approx \nabla J(\mathbf{w})} + \eta \mathbf{v}$$

$$\mathbf{w} := \mathbf{w} + \mathbf{v}$$

where $L^{(i)}(\mathbf{w})$ is based on one observation i and $1 \leq s \leq m$, where m denotes the total number of observations in the data set.



The hyperparameters of the algorithm are

- s - mini-batch size
- α - learning rate
- η - *momentum*, a number between 0 and 1

Momentum Optimization

Example: Path in (w_1, w_2) plane.

Left: no momentum, i.e. $\eta = 0$. Right: Momentum optimization with $\eta > 0$.



Momentum Optimization

Example: Momentum Optimization with $s = 128$, $\alpha = 0.05$, and $\eta = 0.9$.

```
model = models.Sequential()
model.add(layers.Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(layers.Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(layers.Dense(10, activation='softmax'))

model.summary()
```

Model: "sequential_14"

Layer (type)	Output Shape	Param #
dense_35 (Dense)	(None, 512)	401920
dropout_1 (Dropout)	(None, 512)	0
dense_36 (Dense)	(None, 512)	262656
dropout_2 (Dropout)	(None, 512)	0
dense_37 (Dense)	(None, 10)	5130
Total params: 669,706		
Trainable params: 669,706		
Non-trainable params: 0		

```
nepochs = 35
model.compile(loss='categorical_crossentropy', metrics=['accuracy'],
              optimizer=keras.optimizers.SGD(lr=0.05, momentum=0.9))

history = model.fit(X_train, y_train,
                    batch_size=128, epochs=nepochs,
                    verbose=1,
                    validation_data=(X_test, y_test))
```

Nesterov Accelerated Gradient (NAG)

The Nesterov Accelerated Gradient (NAG) algorithm is defined as follows:

First, *momentum vector* \mathbf{v} is initialized at $\mathbf{0}$ and then the updates are

$$\mathbf{v} := -\alpha \underbrace{\frac{1}{s} \sum_{i=1}^s \nabla L^{(i)}(\mathbf{w} + \eta \mathbf{v})}_{\approx \nabla J(\mathbf{w})} + \eta \mathbf{v}$$

$$\mathbf{w} := \mathbf{w} + \mathbf{v}$$

where $L^{(i)}(\mathbf{w} + \eta \mathbf{v})$ is based on one observation i and $1 \leq s \leq m$, where m denotes the total number of observations in the data set.

The hyperparameters of the algorithm are

- s - mini-batch size
- α - learning rate
- η - *momentum*, a number between 0 and 1

Nesterov Accelerated Gradient (NAG)

Example: Nesterov Accelerated Gradient (NAG) with $s = 128$, $\alpha = 0.05$, and $\eta = 0.9$.

```
model = models.Sequential()
model.add(layers.Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(layers.Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(layers.Dense(10, activation='softmax'))
```

```
model.summary()
```

Model: "sequential_14"

Layer (type)	Output Shape	Param #
dense_35 (Dense)	(None, 512)	401920
dropout_1 (Dropout)	(None, 512)	0
dense_36 (Dense)	(None, 512)	262656
dropout_2 (Dropout)	(None, 512)	0
dense_37 (Dense)	(None, 10)	5130
Total params: 669,706		
Trainable params: 669,706		
Non-trainable params: 0		

```
nepochs = 35
model.compile(loss='categorical_crossentropy', metrics=['accuracy'],
              optimizer=keras.optimizers.SGD(lr=0.05, momentum=0.9, nesterov=True))

history = model.fit(X_train, y_train,
                    batch_size=128, epochs=nepochs,
                    verbose=1,
                    validation_data=(X_test, y_test))
```

AdaGrad

The AdaGrad algorithm is defined as follows:

First, initialize vector \mathbf{r} (with $r_k > 0$) and then the updates are

$$\begin{aligned}\mathbf{g} &:= \frac{1}{s} \underbrace{\sum_{i=1}^s \nabla L^{(i)}(\mathbf{w})}_{\approx \nabla J(\mathbf{w})} \\ \mathbf{r} &:= \mathbf{r} + \mathbf{g} \odot \mathbf{g} \\ \mathbf{w} &:= \mathbf{w} - \frac{\alpha}{\sqrt{\mathbf{r} + \epsilon}} \odot \mathbf{g}\end{aligned}$$

where $L^{(i)}(\mathbf{w})$ is based on one observation i and $1 \leq s \leq m$, where m denotes the total number of observations in the data set. \odot denotes element-wise multiplication. The hyperparameters of the algorithm are

- s - mini-batch size
- α - learning rate
- ϵ - positive small parameter, typically around 10^{-7}

AdaGrad

Example: AdaGrad with $s = 128$, $\alpha = 0.05$, $\epsilon = 10^{-5}$, and r_k initialized at 0.1.

```
model = models.Sequential()
model.add(layers.Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(layers.Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(layers.Dense(10, activation='softmax'))

model.summary()
```

Model: "sequential_14"

Layer (type)	Output Shape	Param #
dense_35 (Dense)	(None, 512)	401920
dropout_1 (Dropout)	(None, 512)	0
dense_36 (Dense)	(None, 512)	262656
dropout_2 (Dropout)	(None, 512)	0
dense_37 (Dense)	(None, 10)	5130

=====
Total params: 669,706
Trainable params: 669,706
Non-trainable params: 0

```
nepochs = 3
model.compile(loss='categorical_crossentropy', metrics=['accuracy'],
              optimizer=keras.optimizers.Adagrad(lr=0.05, epsilon=1e-5))

history = model.fit(X_train, y_train,
                    batch_size=128, epochs=nepochs,
                    verbose=1,
                    validation_data=(X_test, y_test))
```

RMSProp

The RMSProp is defined as follows:

First, initialize vector \mathbf{r} (with $r_k > 0$) and then the updates are

$$\begin{aligned}\mathbf{g} &:= \underbrace{\frac{1}{s} \sum_{i=1}^s \nabla L^{(i)}(\mathbf{w})}_{\approx \nabla J(\mathbf{w})} \\ \mathbf{r} &:= \rho \mathbf{r} + (1 - \rho) \mathbf{g} \odot \mathbf{g} \\ \mathbf{w} &:= \mathbf{w} - \frac{\alpha}{\sqrt{\mathbf{r} + \epsilon}} \odot \mathbf{g}\end{aligned}$$

where $L^{(i)}(\mathbf{w})$ is based on one observation i and $1 \leq s \leq m$, where m denotes the total number of observations in the data set. \odot denotes element-wise multiplication. The hyperparameters of the algorithm are

- s - mini-batch size
- α - learning rate
- ϵ - positive small parameter, typically around 10^{-7}
- ρ - decay rate between 0 and 1, typically around 0.9

RMSProp

Example: AdaGrad with $s = 128$, $\alpha = 0.05$, $\epsilon = 10^{-5}$, and $\rho = 0.9$.

```
model = models.Sequential()
model.add(layers.Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(layers.Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(layers.Dense(10, activation='softmax'))

model.summary()
```

Model: "sequential_14"

Layer (type)	Output Shape	Param #
dense_35 (Dense)	(None, 512)	401920
dropout_1 (Dropout)	(None, 512)	0
dense_36 (Dense)	(None, 512)	262656
dropout_2 (Dropout)	(None, 512)	0
dense_37 (Dense)	(None, 10)	5130

=====
Total params: 669,706
Trainable params: 669,706
Non-trainable params: 0

```
nepochs = 35
model.compile(loss='categorical_crossentropy', metrics=['accuracy'],
              optimizer=keras.optimizers.RMSprop(lr=0.05, rho=0.9, epsilon=1e-07))

history = model.fit(X_train, y_train,
                    batch_size=128, epochs=nepochs,
                    verbose=1,
                    validation_data=(X_test, y_test))
```

Adam

Let's combine the Momentum Optimization and RMSProp:

Momentum Optimization	RMSProp
$\mathbf{g} := \frac{1}{s} \underbrace{\sum_{i=1}^s \nabla L^{(i)}(\mathbf{w})}_{\approx \nabla J(\mathbf{w})}$ $\mathbf{v} := \mathbf{g} + \eta \mathbf{v}$ $\mathbf{w} := \mathbf{w} - \alpha \mathbf{v}$	$\mathbf{g} := \frac{1}{s} \underbrace{\sum_{i=1}^s \nabla L^{(i)}(\mathbf{w})}_{\approx \nabla J(\mathbf{w})}$ $\mathbf{r} := \rho \mathbf{r} + (1 - \rho) \mathbf{g} \odot \mathbf{g}$ $\mathbf{w} := \mathbf{w} - \frac{\alpha}{\sqrt{\mathbf{r} + \epsilon}} \odot \mathbf{g}$

to get Adam (Adaptive Momentum):

Adam

First, initialize *momentum vector* $\mathbf{v} = \mathbf{0}$ and vector \mathbf{r} (with $r_k > 0$), then the updates at iteration step t are

$$\mathbf{g} := \underbrace{\frac{1}{s} \sum_{i=1}^s \nabla L^{(i)}(\mathbf{w})}_{\approx \nabla J(\mathbf{w})},$$
$$\mathbf{v} := (1 - \beta_1)\mathbf{g} + \beta_1\mathbf{v}, \quad \mathbf{v} := \frac{\mathbf{v}}{1 - \beta_1^t},$$
$$\mathbf{r} := \beta_2 \mathbf{r} + (1 - \beta_2)\mathbf{g} \odot \mathbf{g}, \quad \mathbf{r} := \frac{\mathbf{r}}{1 - \beta_2^t},$$
$$\mathbf{w} := \mathbf{w} - \frac{\alpha}{\sqrt{\mathbf{r} + \epsilon}} \odot \mathbf{v},$$

where $L^{(i)}(\mathbf{w})$ is based on one observation i and $1 \leq s \leq m$, where m denotes the total number of observations in the data set. \odot denotes element-wise multiplication. The hyperparameters of the algorithm are

- s is the mini-batch size and α is learning rate
- β_1 - *momentum*, a number between 0 and 1 (analogous to η in Momentum Opt.)
- β_2 - decay rate between 0 and 1, typically around 0.9 (analogous to ρ in RMSProp)
- ϵ - positive small parameter, typically around 10^{-7}

Adam

Example: Adam with $s = 128$, $\alpha = 0.001$, $\epsilon = 10^{-7}$, $\beta_1 = 0.9$, and $\beta_2 = 0.999$

```
model = models.Sequential()
model.add(layers.Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(layers.Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(layers.Dense(10, activation='softmax'))

model.summary()
```

Model: "sequential_14"

Layer (type)	Output Shape	Param #
dense_35 (Dense)	(None, 512)	401920
dropout_1 (Dropout)	(None, 512)	0
dense_36 (Dense)	(None, 512)	262656
dropout_2 (Dropout)	(None, 512)	0
dense_37 (Dense)	(None, 10)	5130
Total params: 669,706		
Trainable params: 669,706		
Non-trainable params: 0		

```
nepochs = 35
model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='adam')

history = model.fit(X_train, y_train,
                    batch_size=128, epochs=nepochs,
                    verbose=1,
                    validation_data=(X_test, y_test))
```


Adam

Example: Adam with $s = 128$, $\alpha = 0.05$, $\epsilon = 10^{-5}$, $\beta_1 = 0.85$, and $\beta_2 = 0.95$

```
model = models.Sequential()
model.add(layers.Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(layers.Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(layers.Dense(10, activation='softmax'))

model.summary()
```

Model: "sequential_14"

Layer (type)	Output Shape	Param #
dense_35 (Dense)	(None, 512)	401920
dropout_1 (Dropout)	(None, 512)	0
dense_36 (Dense)	(None, 512)	262656
dropout_2 (Dropout)	(None, 512)	0
dense_37 (Dense)	(None, 10)	5130

=====

Total params: 669,706
Trainable params: 669,706
Non-trainable params: 0

```
nepochs = 35

model.compile(loss='categorical_crossentropy', metrics=['accuracy'],
              optimizer=keras.optimizers.adam(lr=0.05, beta_1=0.85, beta_2=0.95, epsilon=1e-05))

history = model.fit(X_train, y_train,
                    batch_size=128, epochs=nepochs,
                    verbose=1,
                    validation_data=(X_test, y_test))
```